

madrid_2009

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, LogisticRegression, Lasso, Ridge
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2009\madrid_2009.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM
0	2009-10-01 01:00:00	NaN	0.27	NaN	NaN	NaN	39.889999	48.150002	NaN	50.680000	18.2600
1	2009-10-01 01:00:00	NaN	0.22	NaN	NaN	NaN	21.230000	24.260000	NaN	55.880001	10.5800
2	2009-10-01 01:00:00	NaN	0.18	NaN	NaN	NaN	31.230000	34.880001	NaN	49.060001	25.1900
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.5300
4	2009-10-01 01:00:00	NaN	0.41	NaN	NaN	0.12	61.349998	76.260002	NaN	38.090000	23.7600
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.8300
215684	2009-06-01 00:00:00	NaN	0.31	NaN	NaN	NaN	76.110001	101.099998	NaN	41.220001	9.9200
215685	2009-06-01 00:00:00	0.13	NaN	0.86	NaN	0.23	81.050003	99.849998	NaN	24.830000	12.4600
215686	2009-06-01 00:00:00	0.21	NaN	2.96	NaN	0.10	72.419998	82.959999	NaN	NaN	13.0300
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.3600

215688 rows × 12 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 215688 entries, 0 to 215687
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        215688 non-null  object
1   BEN         60082 non-null   float64
2   CO          190801 non-null  float64
3   EBE         60081 non-null   float64
4   MXY         24846 non-null   float64
5   NMHC        74748 non-null   float64
6   NO_2        214562 non-null  float64
7   NOx         214565 non-null  float64
8   OXY         24854 non-null   float64
9   O_3         204482 non-null  float64
10  PM10        196331 non-null  float64
11  PM25        55822 non-null   float64
12  PXY         24854 non-null   float64
13  SO_2        212671 non-null  float64
14  TCH         75213 non-null   float64
15  TOL         59920 non-null   float64
16  station     215688 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 28.0+ MB
```

```
In [4]: df1=df.dropna()  
df1
```

Out[4]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM1
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.53000
20	2009-10-01 01:00:00	0.38	0.32	0.32	0.89	0.01	17.969999	19.240000	1.00	65.870003	10.52000
24	2009-10-01 01:00:00	0.55	0.24	0.65	1.79	0.18	36.619999	43.919998	1.28	48.070000	19.15000
28	2009-10-01 02:00:00	0.65	0.21	1.20	2.04	0.18	37.169998	48.869999	1.21	26.950001	32.20000
45	2009-10-01 02:00:00	0.38	0.30	0.50	1.15	0.00	17.889999	19.299999	1.00	60.009998	12.26000
...
215659	2009-05-31 23:00:00	0.54	0.27	1.00	0.69	0.09	28.280001	29.490000	0.86	78.750000	15.17000
215663	2009-05-31 23:00:00	0.74	0.35	1.13	1.65	0.15	56.410000	69.870003	1.26	56.799999	11.80000
215667	2009-06-01 00:00:00	0.78	0.29	0.99	1.96	0.04	64.870003	82.629997	1.13	58.000000	12.67000
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.83000
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.36000


24717 rows × 17 columns



```
In [5]: df1=df1.drop(["date"],axis=1)
```


```
In [6]: sns.heatmap(df1.corr())
```

```
Out[6]: <AxesSubplot:>
```



```
In [7]: plt.plot(df1["EBE"],df1["PXY"],"o")
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x1b59f30f250>]
```



```
In [8]: data=df[["EBE","PXY"]]
```

```
In [9]: # sns.stripplot(x=df["EBE"],y=df["PXY"],jitter=True,marker='o',color='blue')
```

```
In [10]: x=df1.drop(["EBE"],axis=1)  
y=df1["EBE"]  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


LINEAR

```
In [11]: li=LinearRegression()  
li.fit(x_train,y_train)
```

```
Out[11]: LinearRegression()
```

```
In [12]: prediction=li.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[12]: <matplotlib.collections.PathCollection at 0x1b59f3c8f40>
```



```
In [13]: lis=li.score(x_test,y_test)
```

```
In [14]: df1["TCH"].value_counts()
```

```
Out[14]: 1.39    1091  
1.36    1056  
1.38    1046  
1.40    1018  
1.37    1017  
...  
2.52      1  
1.16      1  
2.41      1  
1.13      1  
2.79      1  
Name: TCH, Length: 169, dtype: int64
```

```
In [15]: df1.loc[df1["TCH"]<1.40,"TCH"]=1  
df1.loc[df1["TCH"]>1.40,"TCH"]=2  
df1["TCH"].value_counts()
```

```
Out[15]: 1.0    12963  
2.0    11754  
Name: TCH, dtype: int64
```


```
In [16]: # Lasso
```

```
In [17]: la=Lasso(alpha=5)
la.fit(x_train,y_train)
```

```
Out[17]: Lasso(alpha=5)
```

```
In [18]: prediction1=la.predict(x_test)
plt.scatter(y_test,prediction1)
```

```
Out[18]: <matplotlib.collections.PathCollection at 0x1b59ffde310>
```



```
In [19]: las=la.score(x_test,y_test)
```


RIDGE

```
In [20]: rr=Ridge(alpha=1)
rr.fit(x_train,y_train)
```

```
Out[20]: Ridge(alpha=1)
```

```
In [21]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

Out[21]: <matplotlib.collections.PathCollection at 0x1b59f3326a0>



```
In [22]: rrs=rr.score(x_test,y_test)
```


ElasticNet

```
In [23]: en=ElasticNet()
en.fit(x_train,y_train)
```

Out[23]: ElasticNet()

```
In [24]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

Out[24]: <matplotlib.collections.PathCollection at 0x1b5a005e2b0>



```
In [25]: ens=en.score(x_test,y_test)
```

```
In [26]: print(rr.score(x_test,y_test))  
rr.score(x_train,y_train)
```

0.8655545107640262

Out[26]: 0.891088111281474

LOGISTIC

```
In [27]: g={"TCH":{1.0:"Low",2.0:"High"}}  
df1=df1.replace(g)  
df1["TCH"].value_counts()
```

Out[27]: Low 12963
High 11754
Name: TCH, dtype: int64

```
In [28]: x=df1.drop(["TCH"],axis=1)  
y=df1["TCH"]  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [29]: lo=LogisticRegression()  
lo.fit(x_train,y_train)
```

Out[29]: LogisticRegression()

```
In [30]: prediction3=lo.predict(x_test)  
plt.scatter(y_test,prediction3)
```

Out[30]: <matplotlib.collections.PathCollection at 0x1b59fd047c0>



```
In [31]: los=lo.score(x_test,y_test)
```


Random Forest

```
In [32]: from sklearn.ensemble import RandomForestClassifier
        from sklearn.model_selection import GridSearchCV
```

```
In [33]: g1={"TCH":{"Low":1.0,"High":2.0}}
        df1=df1.replace(g1)
```

```
In [34]: x=df1.drop(["TCH"],axis=1)
        y=df1["TCH"]
        x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [35]: rfc=RandomForestClassifier()
        rfc.fit(x_train,y_train)
```

```
Out[35]: RandomForestClassifier()
```

```
In [36]: parameter={
        'max_depth':[1,2,4,5,6],
        'min_samples_leaf':[5,10,15,20,25],
        'n_estimators':[10,20,30,40,50]
        }
```

```
In [37]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
        grid_search.fit(x_train,y_train)
```

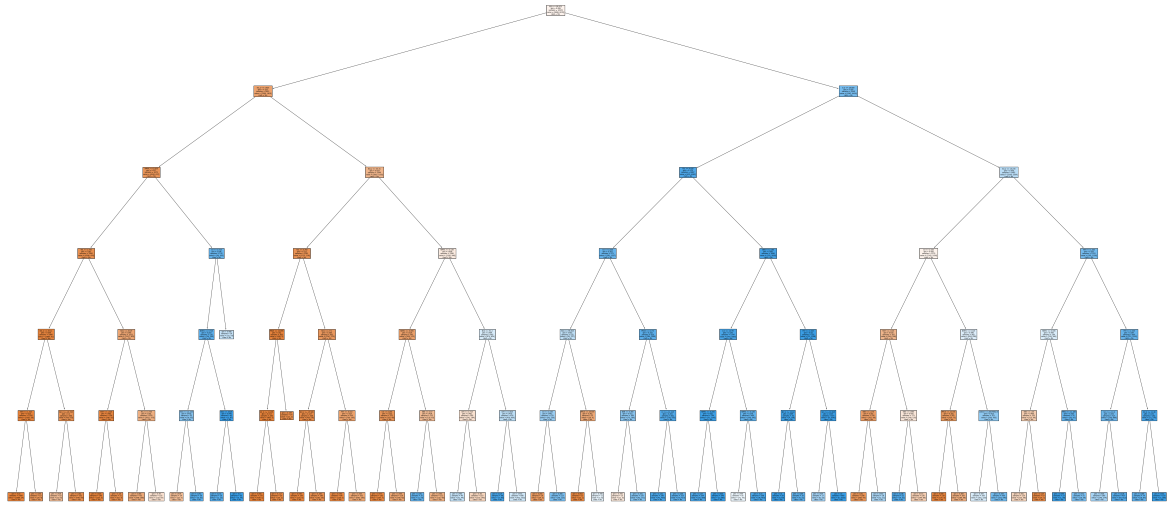
```
Out[37]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
        param_grid={'max_depth': [1, 2, 4, 5, 6],
        'min_samples_leaf': [5, 10, 15, 20, 25],
        'n_estimators': [10, 20, 30, 40, 50]},
        scoring='accuracy')
```

```
In [38]: rfcs=grid_search.best_score_
```

```
In [39]: rfc_best=grid_search.best_estimator_
```

```
In [40]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],
          Text(4425.517241379311, 155.3142857142857, 'gini = 0.031\nsamples = 118\nvalue = [3, 189]\nclass = No'))
```



```
In [41]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.8655713958460309
Lasso: 0.3459256380275074
Ridge: 0.8655545107640262
ElasticNet: 0.5761325583219384
Logistic: 0.5230582524271845
Random Forest: 0.8623202837321089
```

Best Model is Random Forest

madrid_2010

```
In [42]: df2=pd.read_csv(r"C:\Users\user\Downloads\csvs_per_year\csvs_per_year\madrid_2010\df2")
```

Out[42]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PI
0	2010-03-01 01:00:00	NaN	0.29	NaN	NaN	NaN	25.090000	29.219999	NaN	68.930000	I
1	2010-03-01 01:00:00	NaN	0.27	NaN	NaN	NaN	24.879999	30.040001	NaN	NaN	I
2	2010-03-01 01:00:00	NaN	0.28	NaN	NaN	NaN	17.410000	20.540001	NaN	72.120003	I
3	2010-03-01 01:00:00	0.38	0.24	1.74	NaN	0.05	15.610000	21.080000	NaN	72.970001	19.410
4	2010-03-01 01:00:00	0.79	NaN	1.32	NaN	NaN	21.430000	26.070000	NaN	NaN	24.670
...
209443	2010-08-01 00:00:00	NaN	0.55	NaN	NaN	NaN	125.000000	219.899994	NaN	25.379999	I
209444	2010-08-01 00:00:00	NaN	0.27	NaN	NaN	NaN	45.709999	47.410000	NaN	NaN	51.259
209445	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	0.24	46.560001	49.040001	NaN	46.250000	I
209446	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	46.770000	50.119999	NaN	77.709999	I
209447	2010-08-01 00:00:00	0.92	0.43	0.71	NaN	0.25	76.330002	88.190002	NaN	52.259998	47.150

209448 rows × 17 columns



```
In [43]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209448 entries, 0 to 209447
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209448 non-null  object
1   BEN         60268 non-null   float64
2   CO          94982 non-null   float64
3   EBE         60253 non-null   float64
4   MXY         6750 non-null    float64
5   NMHC        51727 non-null   float64
6   NO_2        208219 non-null   float64
7   NOx         208210 non-null   float64
8   OXY         6750 non-null    float64
9   O_3         126684 non-null   float64
10  PM10        106186 non-null   float64
11  PM25        55514 non-null   float64
12  PXY         6740 non-null    float64
13  SO_2        93184 non-null   float64
14  TCH         51730 non-null   float64
15  TOL         60171 non-null   float64
16  station     209448 non-null   int64
dtypes: float64(15), int64(1), object(1)
memory usage: 27.2+ MB
```

```
In [44]: df3=df2.dropna()  
df3
```

Out[44]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM1
11	2010-03-01 01:00:00	0.78	0.18	0.84	0.73	0.28	10.420000	11.900000	1.0	90.309998	18.37000
23	2010-03-01 01:00:00	0.70	0.23	1.00	0.73	0.18	17.820000	22.290001	1.0	70.550003	23.63999
35	2010-03-01 02:00:00	0.58	0.17	0.84	0.73	0.28	3.500000	4.950000	1.0	68.849998	5.60000
47	2010-03-01 02:00:00	0.33	0.21	0.84	0.73	0.17	10.810000	14.900000	1.0	74.750000	7.89000
59	2010-03-01 03:00:00	0.38	0.16	0.64	1.00	0.26	2.750000	4.200000	1.0	93.629997	5.13000
...
191879	2010-05-31 22:00:00	0.60	0.26	0.82	0.13	0.16	33.360001	43.779999	1.0	38.459999	20.34000
191891	2010-05-31 23:00:00	0.41	0.16	0.71	0.19	0.10	24.299999	26.059999	1.0	50.290001	14.38000
191903	2010-05-31 23:00:00	0.57	0.28	0.64	0.19	0.18	35.540001	44.590000	1.0	34.020000	22.84000
191915	2010-06-01 00:00:00	0.34	0.16	0.69	0.22	0.10	23.559999	25.209999	1.0	45.930000	10.77000
191927	2010-06-01 00:00:00	0.43	0.25	0.79	0.22	0.18	34.910000	42.369999	1.0	29.540001	15.35000


6666 rows × 17 columns



```
In [45]: df3=df3.drop(["date"],axis=1)
```

```
In [46]: sns.heatmap(df3.corr())
```

```
Out[46]: <AxesSubplot:>
```



```
In [47]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


Linear

```
In [48]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[48]: LinearRegression()
```

```
In [49]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[49]: <matplotlib.collections.PathCollection at 0x1b5a00dc4c0>
```



```
In [50]: lis=li.score(x_test,y_test)
```

```
In [51]: df3["TCH"].value_counts()
```

```
Out[51]: 1.36    364
         1.38    351
         1.39    324
         1.35    323
         1.37    321
         ...
         2.07     1
         2.17     1
         2.53     1
         2.12     1
         2.05     1
         Name: TCH, Length: 100, dtype: int64
```

```
In [52]: df3.loc[df3["TCH"]<1.40,"TCH"]=1
         df3.loc[df3["TCH"]>1.40,"TCH"]=2
         df3["TCH"].value_counts()
```

```
Out[52]: 1.0    3340
         2.0    3326
         Name: TCH, dtype: int64
```


Lasso

```
In [53]: la=Lasso(alpha=5)
         la.fit(x_train,y_train)
```

```
Out[53]: Lasso(alpha=5)
```

```
In [54]: prediction1=la.predict(x_test)
         plt.scatter(y_test,prediction1)
```

```
Out[54]: <matplotlib.collections.PathCollection at 0x1b5a01384c0>
```



```
In [55]: las=la.score(x_test,y_test)
```


Ridge

```
In [56]: rr=Ridge(alpha=1)  
rr.fit(x_train,y_train)
```

```
Out[56]: Ridge(alpha=1)
```

```
In [57]: prediction2=rr.predict(x_test)  
plt.scatter(y_test,prediction2)
```

```
Out[57]: <matplotlib.collections.PathCollection at 0x1b5a01974c0>
```



```
In [58]: rrs=rr.score(x_test,y_test)
```

ElasticNet


```
In [59]: en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[59]: ElasticNet()
```



```
In [60]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[60]: <matplotlib.collections.PathCollection at 0x1b5a01ebf40>
```



```
In [61]: ens=en.score(x_test,y_test)
```

```
In [62]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.4588987162157072
```

```
Out[62]: 0.44755643852232574
```

Logistic

```
In [63]: g={"TCH":{1.0:"Low",2.0:"High"}}
df3=df3.replace(g)
df3["TCH"].value_counts()
```

```
Out[63]: Low      3340
High      3326
Name: TCH, dtype: int64
```


```
In [64]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [65]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[65]: LogisticRegression()
```

```
In [66]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[66]: <matplotlib.collections.PathCollection at 0x1b59fbd1ee0>
```



```
In [67]: los=lo.score(x_test,y_test)
```

Random Forest

```
In [68]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [69]: g1={"TCH":{"Low":1.0,"High":2.0}}
df3=df3.replace(g1)
```

```
In [70]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [71]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[71]: RandomForestClassifier()
```

```
In [72]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [73]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[73]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 4, 5, 6],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [74]: rfcs=grid_search.best_score_
```

```
In [75]: rfc_best=grid_search.best_estimator_
```

```
In [76]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'])
```

```
Out[76]: [Text(2080.734939759036, 2019.0857142857144, 'BEN <= 0.735\nngini = 0.5\nsamples = 2926\nvalue = [2353, 2313]\nclass = Yes'),
Text(907.5903614457832, 1708.457142857143, 'O_3 <= 42.015\nngini = 0.464\nsamples = 2045\nvalue = [2081, 1196]\nclass = Yes'),
Text(255.4698795180723, 1397.8285714285716, 'MXY <= 0.185\nngini = 0.473\nsamples = 387\nvalue = [239, 385]\nclass = No'),
Text(107.56626506024097, 1087.2, 'CO <= 0.205\nngini = 0.108\nsamples = 21\nvalue = [33, 2]\nclass = Yes'),
Text(53.78313253012048, 776.5714285714287, 'gini = 0.0\nsamples = 10\nvalue = [18, 0]\nclass = Yes'),
Text(161.34939759036143, 776.5714285714287, 'gini = 0.208\nsamples = 11\nvalue = [15, 2]\nclass = Yes'),
Text(403.3734939759036, 1087.2, 'PM10 <= 16.33\nngini = 0.455\nsamples = 366\nvalue = [206, 383]\nclass = No'),
Text(268.9156626506024, 776.5714285714287, 'NOx <= 33.24\nngini = 0.489\nsamples = 218\nvalue = [150, 201]\nclass = No'),
Text(161.34939759036143, 465.9428571428573, 'PM25 <= 5.665\nngini = 0.477\nsamples = 58\nvalue = [57, 37]\nclass = Yes'),
Text(107.56626506024097, 155.3142857142857, 'gini = 0.26\nsamples = 15\nvalue = [50, 17]\nclass = Yes')]
```

```
In [77]: print("Linear:",lis)
         print("Lasso:",las)
         print("Ridge:",rrs)
         print("ElasticNet:",ens)
         print("Logistic:",los)
         print("Random Forest:",rfcs)
```

```
Linear: 0.45900314075246096
Lasso: -2.087909812176214e-05
Ridge: 0.4588987162157072
ElasticNet: 0.34488908831719434
Logistic: 0.487
Random Forest: 0.7801114444920703
```

```
In [ ]:
```